

# Understanding Job Execution in Supercomputers

Alex Moore, UNC Charlotte  
Di Zhang, High Performance Computing (HPC)



## Introduction

Super Computers are built to solve complex problems that require high-speed data processing, such as:

- Spacecraft simulation
- Analyzing cyber attacks and natural disasters

This is done by running programs or jobs, which are then scheduled through a process known as job scheduling.

Why is job scheduling important?

- Minimizes job waiting time, slowdown, or completion time.
- Allows for more complex problems to be solved.

## Objectives

Methods to improve job scheduling:

- **TRIP** “Trade-off between Prediction Accuracy and Underestimation Rate in Job Runtime Estimates”
  - ML framework that utilizes data censoring to improve prediction accuracy with a low underestimation rate of job runtimes.
- **SchedInspector** “SchedInspector: A Batch Job Scheduling Inspector Using Reinforcement Learning”
  - Determines the competency of scheduled jobs by examining the decisions of the base job scheduler.

## Method

### Phase 1: Learning

During this phase I learned:

- Process of estimating HPC Batch Job Schedulers Runtime
- Methods of improving prediction accuracy and job execution performance
  - TRIP
  - SchedInspector
- Fundamentals of Machine Learning
  - Reinforcement Learning
- Pandas
- Jupyter Notebook

### Phase 2: Data Mapping

- Utilized Jupyter Notebook and Pandas to map data from previous traces
  - Submit time, Wait time, Run time, Requested Time, User ID, etc

### Phase 3: Machine Learning: Feature Extraction

- Extracted features from previous job traces to determine trends amongst the data

TABLE III: The features used for runtime prediction. Here, a class contains jobs with the same user name, project name, and job name.

Feature	Description
$t_{last1}$	the actual runtime of the last job of the same class
$t_{last2}$	the actual runtime of the second-to-last job of the same class
$t_{supplied}$	user-supplied job runtime estimate
$n_{supplied}$	the number of nodes requested by the user
$A_{average}$	the average accuracy of the historical jobs of the same class
$A_{max}$	the maximum accuracy of the historical jobs of the same class
$t_{longest}$	the longest actual runtime of the historical jobs of the same class
$t_{longest10}$	the longest actual runtime of the ten last jobs of the same class
$t_{average}$	the average actual runtime of the historical jobs of the same class
$t_{average10}$	the average actual runtime of ten last jobs of the same class
$t_{percentile25}$	the actual runtime of the 25th percentile historical jobs of the same class

## Results

- Mapped data from previous job traces
- Successfully extracted each of the listed features

	Job Number	Submit Time	Wait Time	Run Time	Allocated Processors	Average CPU time used	Used Memory	Requested Processors
0	11	566129	5	28826	1	27758	-1	1
1	12	566290	532	26171	1	24666	-1	1
2	13	567314	15757	8071	8	7334	-1	8
3	14	571164	21588	64832	32	44134	-1	32
4	15	574284	4647	64384	7	58033	-1	7
...	...	...	...	...	...	...	...	...
59710	73492	63562563	11	71	4	40.25	-1	4
59711	73493	63562705	24	75	4	42.00	-1	4
59712	73494	63563061	11	72	4	40.25	-1	4

Mapping of data from SDSC-SP-1998-4.2 (1)

Requested Time	Requested Memory	Status	User ID	Group ID	Application Number	Queue Number	Partition Number
28800	-1	5	153	75	18180	3	-1
28800	-1	1	153	75	18184	3	-1
64800	-1	1	150	6	13592	4	-1
64800	-1	5	6	6	13602	3	-1
64800	-1	1	151	75	13607	3	-1
...	...	...	...	...	...	...	...
1200	-1	1	3	74	56633	1	-1
1200	-1	1	3	74	56636	1	-1
1200	-1	1	3	74	56638	1	-1

Mapping of data from SDSC-SP-1998-4.2 (2)

```
#Feature tlongest
#the longest actual runtime of the historical jobs of the same class

#groups data by class of job
grouped = df.groupby(["User ID"])

#max value of "Run Time" column for each group
max_runtimes = grouped["Run Time"].max()

#longest actual runtime of historical jobs of same class
longest_runtime = max_runtimes.max()

print(longest_runtime)

9996
```

Example of feature tlongest from SDSC-SP-1998-4.2



Hewlett Packard's Frontier; The world's fastest supercomputer

## Conclusions

- This data provides valuable, real-world insight into the performance of job scheduling.
- Features will be compared with other past job traces to determine trends in the data.
- These trends will help in the implementation of the TRIP model and SchedInspector.
- This will allow for a comparison of their benefits and look further into other methods of improving job scheduling performance.

## References

Yuping Fan, Zhiling Lan, Paul Rich, William E. Allcock, Michael E. Papka Trade-off between Prediction Accuracy and Underestimation Rate in Job Runtime Estimates Table III  
Di Zhang, Dong Dai, Bing Xie SchedInspector: A Batch Job Scheduling Inspector Using Reinforcement Learning